



Get To Know Your Amazon Elastic File System (EFS)

Priya Chakravarthi – Product Management
Nathan Redmond – Solutions Architect

Agenda

- Introduction
- Deep dive on Amazon EFS
- What's new?
- Customer use cases
- Demo
- Summary

How can you transform faster with file storage?



Infrastructure

Improve fundamentals:
security, availability,
performance, and cost



Architecture

Increase agility and
ability to innovate



Business

Maximize your
business results

Accelerating your infrastructure transformation



Infrastructure

- Provide highly secure and reliable storage, always available and accessible
- Deliver price and performance to match your application needs at scale
- Enable native data protection, compliance, and backup

Accelerating your architectural transformation



Architecture

- Capitalize on what the cloud offers by integrating natively with AWS services
- Secure data using native security methods, encryption, and IAM authorization
- Support latest compute models such as containers, microservices, and serverless

Accelerating your business transformation



Business

- Reduce the cost and complexity of managing, maintaining, and operating infrastructure
- Generate opportunities to increase business value through analytics and machine learning
- Enable you to focus on innovation

Using the right tool



**Amazon Elastic
File Service (EFS)**

**Linux-Based
Workloads**



**Amazon FSx
for Windows File Server**

**Windows-Based
Workloads**



**Amazon FSx
for Lustre**

**Compute-
Intensive
Workloads**

FILE SYSTEMS FOR BUSINESS WORKLOADS


**FILE SYSTEM FOR COMPUTE-
INTENSIVE WORKLOADS**

Amazon EFS

Amazon Elastic File System (Amazon EFS)




Amazon EFS
is a fully managed file system that is...




Use cases for Amazon EFS



Home directories
Container storage
Application test/dev



Lift and shift enterprise apps
Web serving
Content management
Database backups




Analytics
Media workflows

Metadata-intensive jobs

Scale-out jobs


Low latency and serial I/O

High throughput and parallel I/O



Elastic

- Grow & shrink on demand
- No need to provision and manage infrastructure & capacity
- Pay as you go, pay only for what you use
- Simple to use, create a file system in seconds



Scalable

- Grow up to petabytes
- Performance modes for low latencies and maximum I/O
- Throughput that scales with storage
- Provisioned throughput available



Integrated

- Integrated with various AWS computing models
- Shared access from on-premises, inter region, and cloud-native applications
- Access concurrently from thousands of Amazon EC2 instances
- Attach to containers launched by both Amazon ECS (AWS Fargate) and Amazon EKS (AWS Fargate - coming soon)
- Use with Amazon SageMaker notebooks



Amazon EFS



Highly available, durable

- Stores data across three availability zones for high availability and durability
- Access your file system from multiple AZs
- Strong consistency for concurrent access



Secure

- Control network traffic
- Control file and directory access
- Control administrative (API) access with AWS IAM
- Encrypt data at rest and in transit
- Control NFS client access with AWS IAM
- Manage application access with EFS Access Points



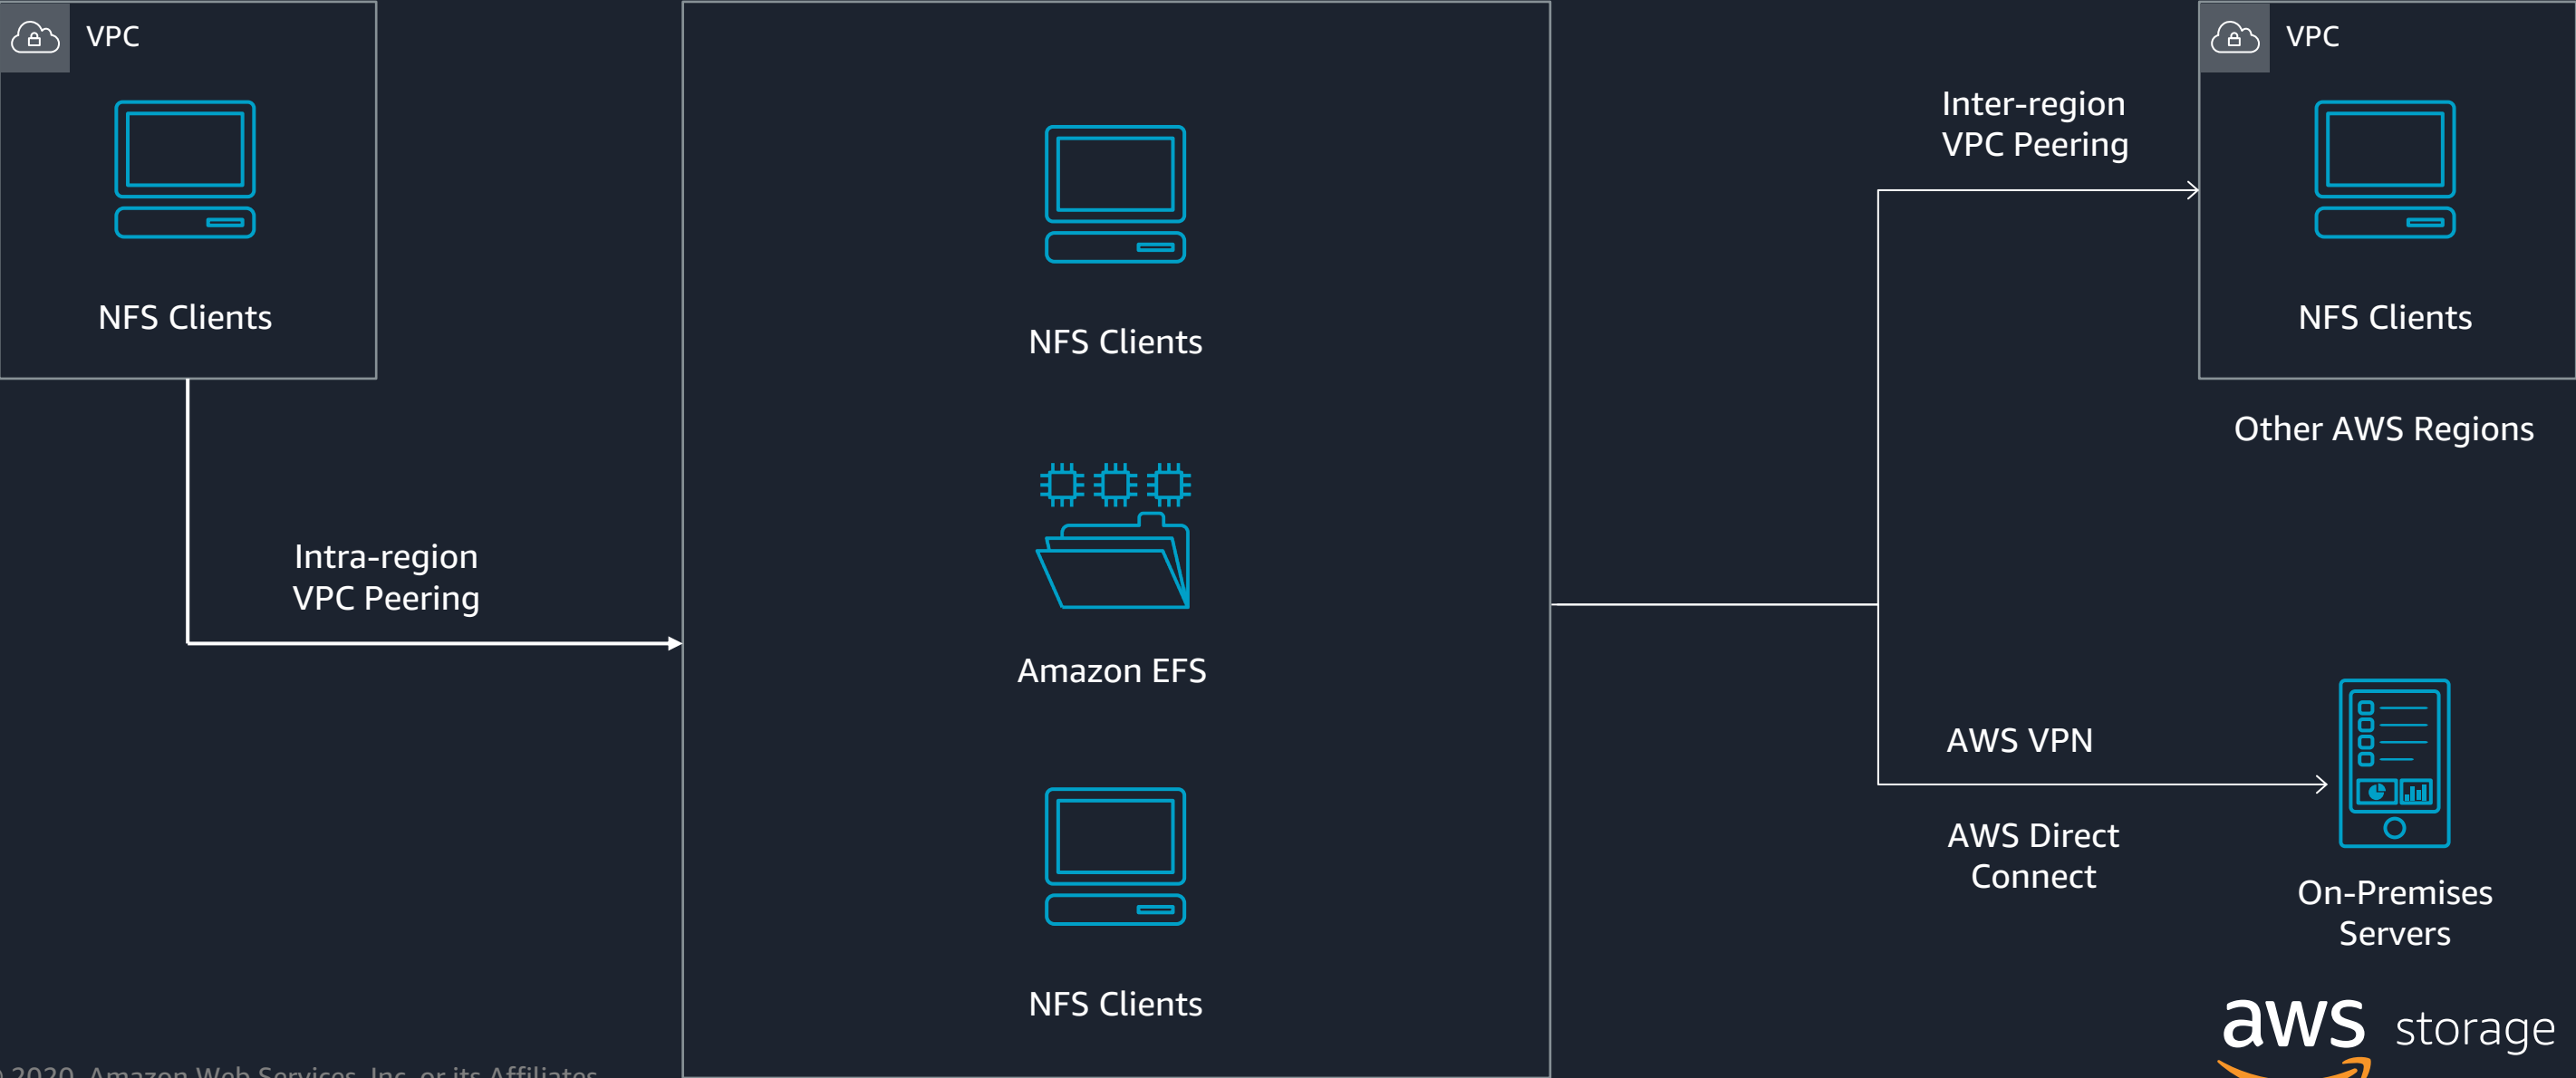
Global footprint

- Amazon EFS is available in ALL regions
- Globally expanding regional service
- 90-day SLA for launch in new regions: Cape Town and Milan (5/13)

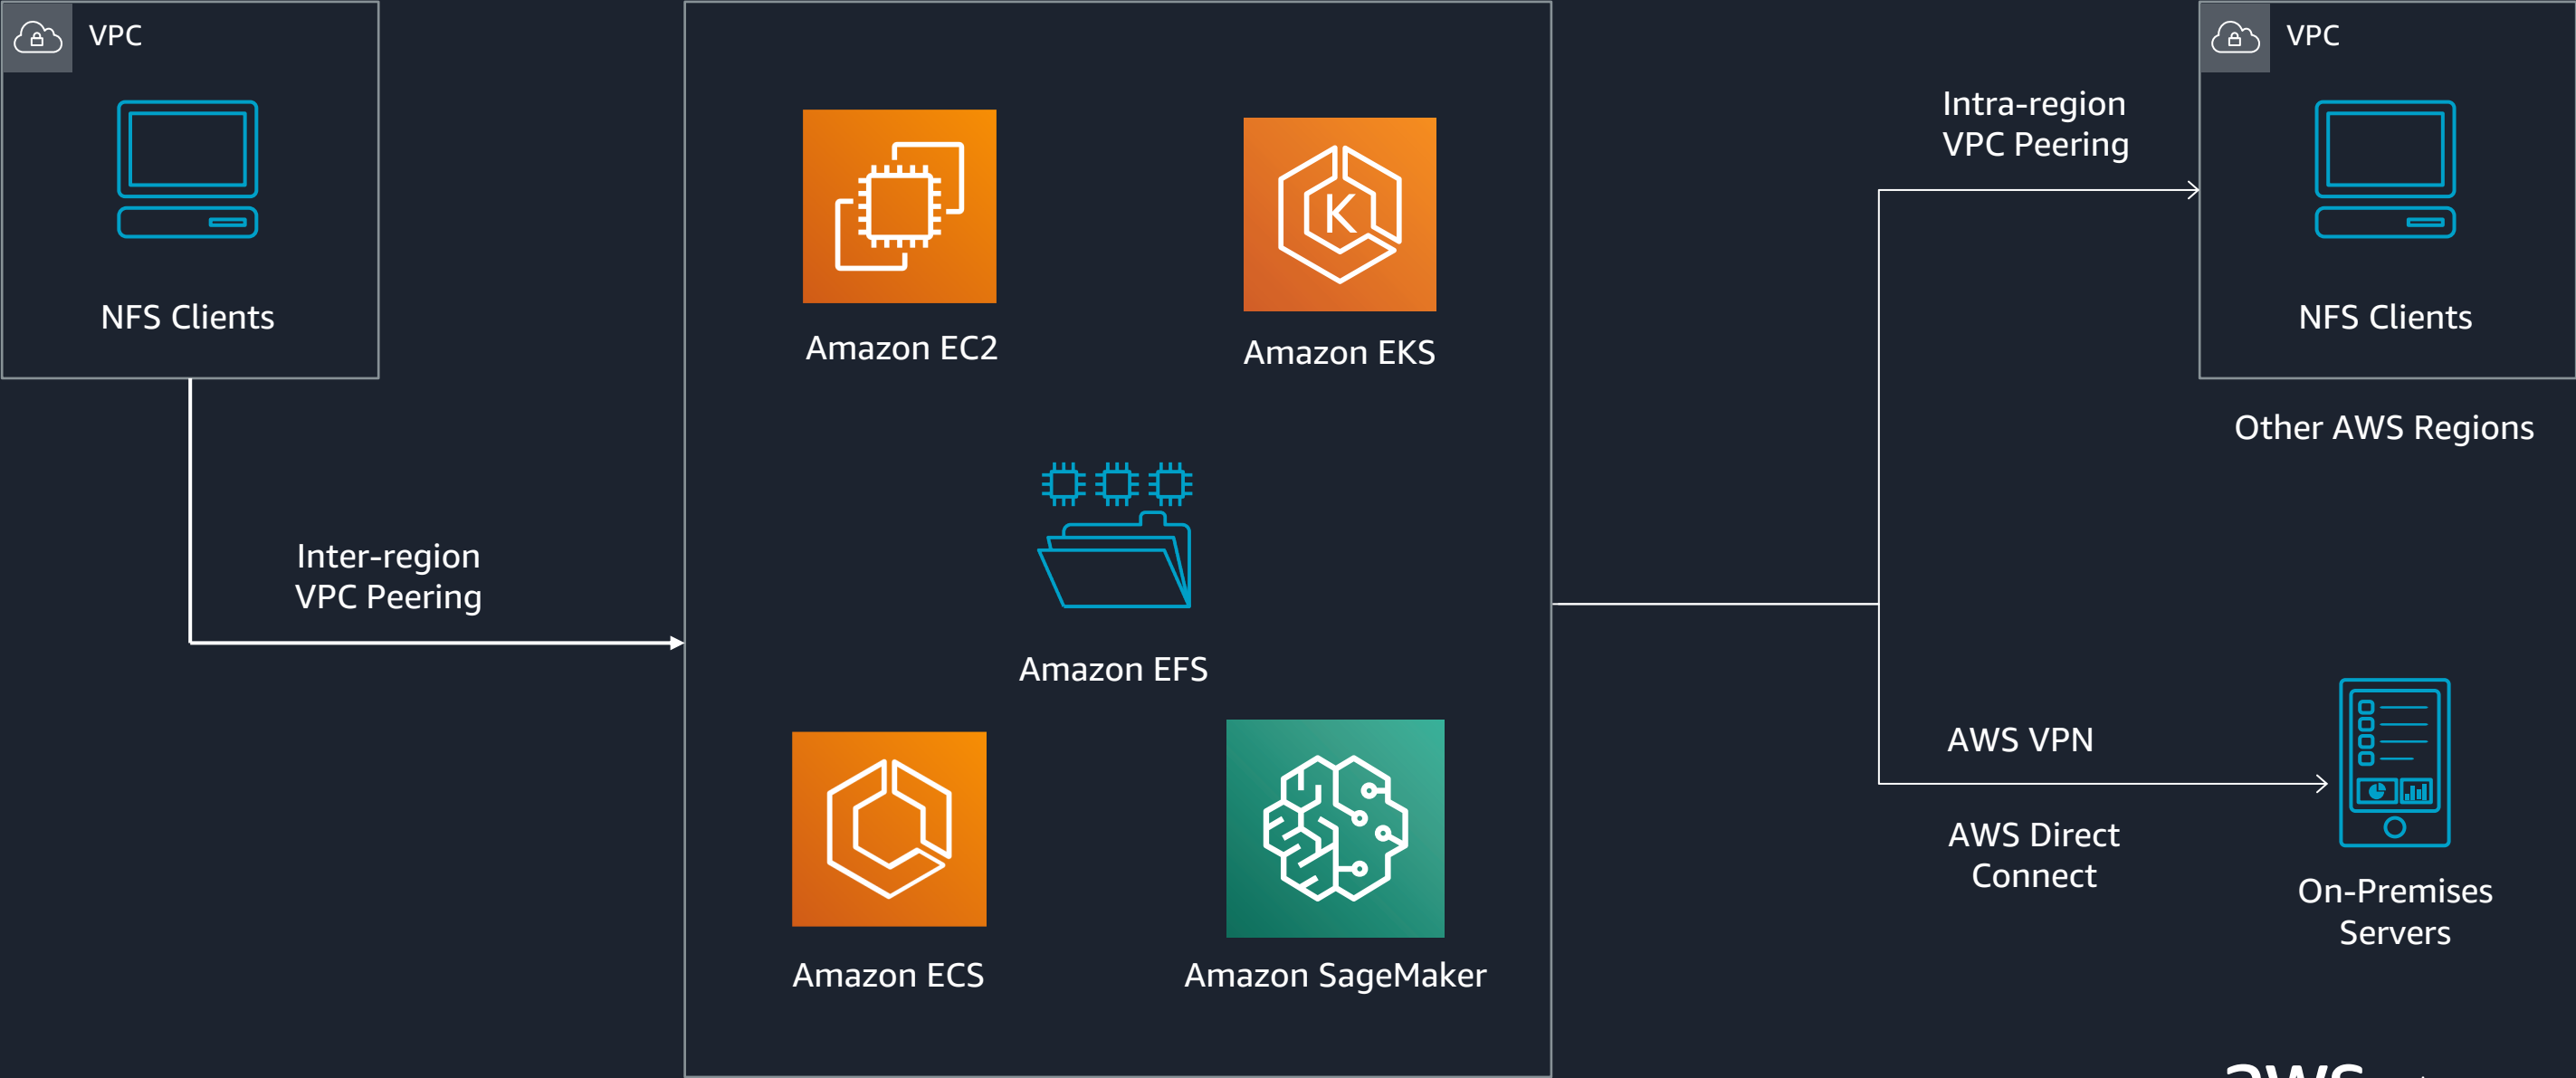


Amazon EFS

Shared access



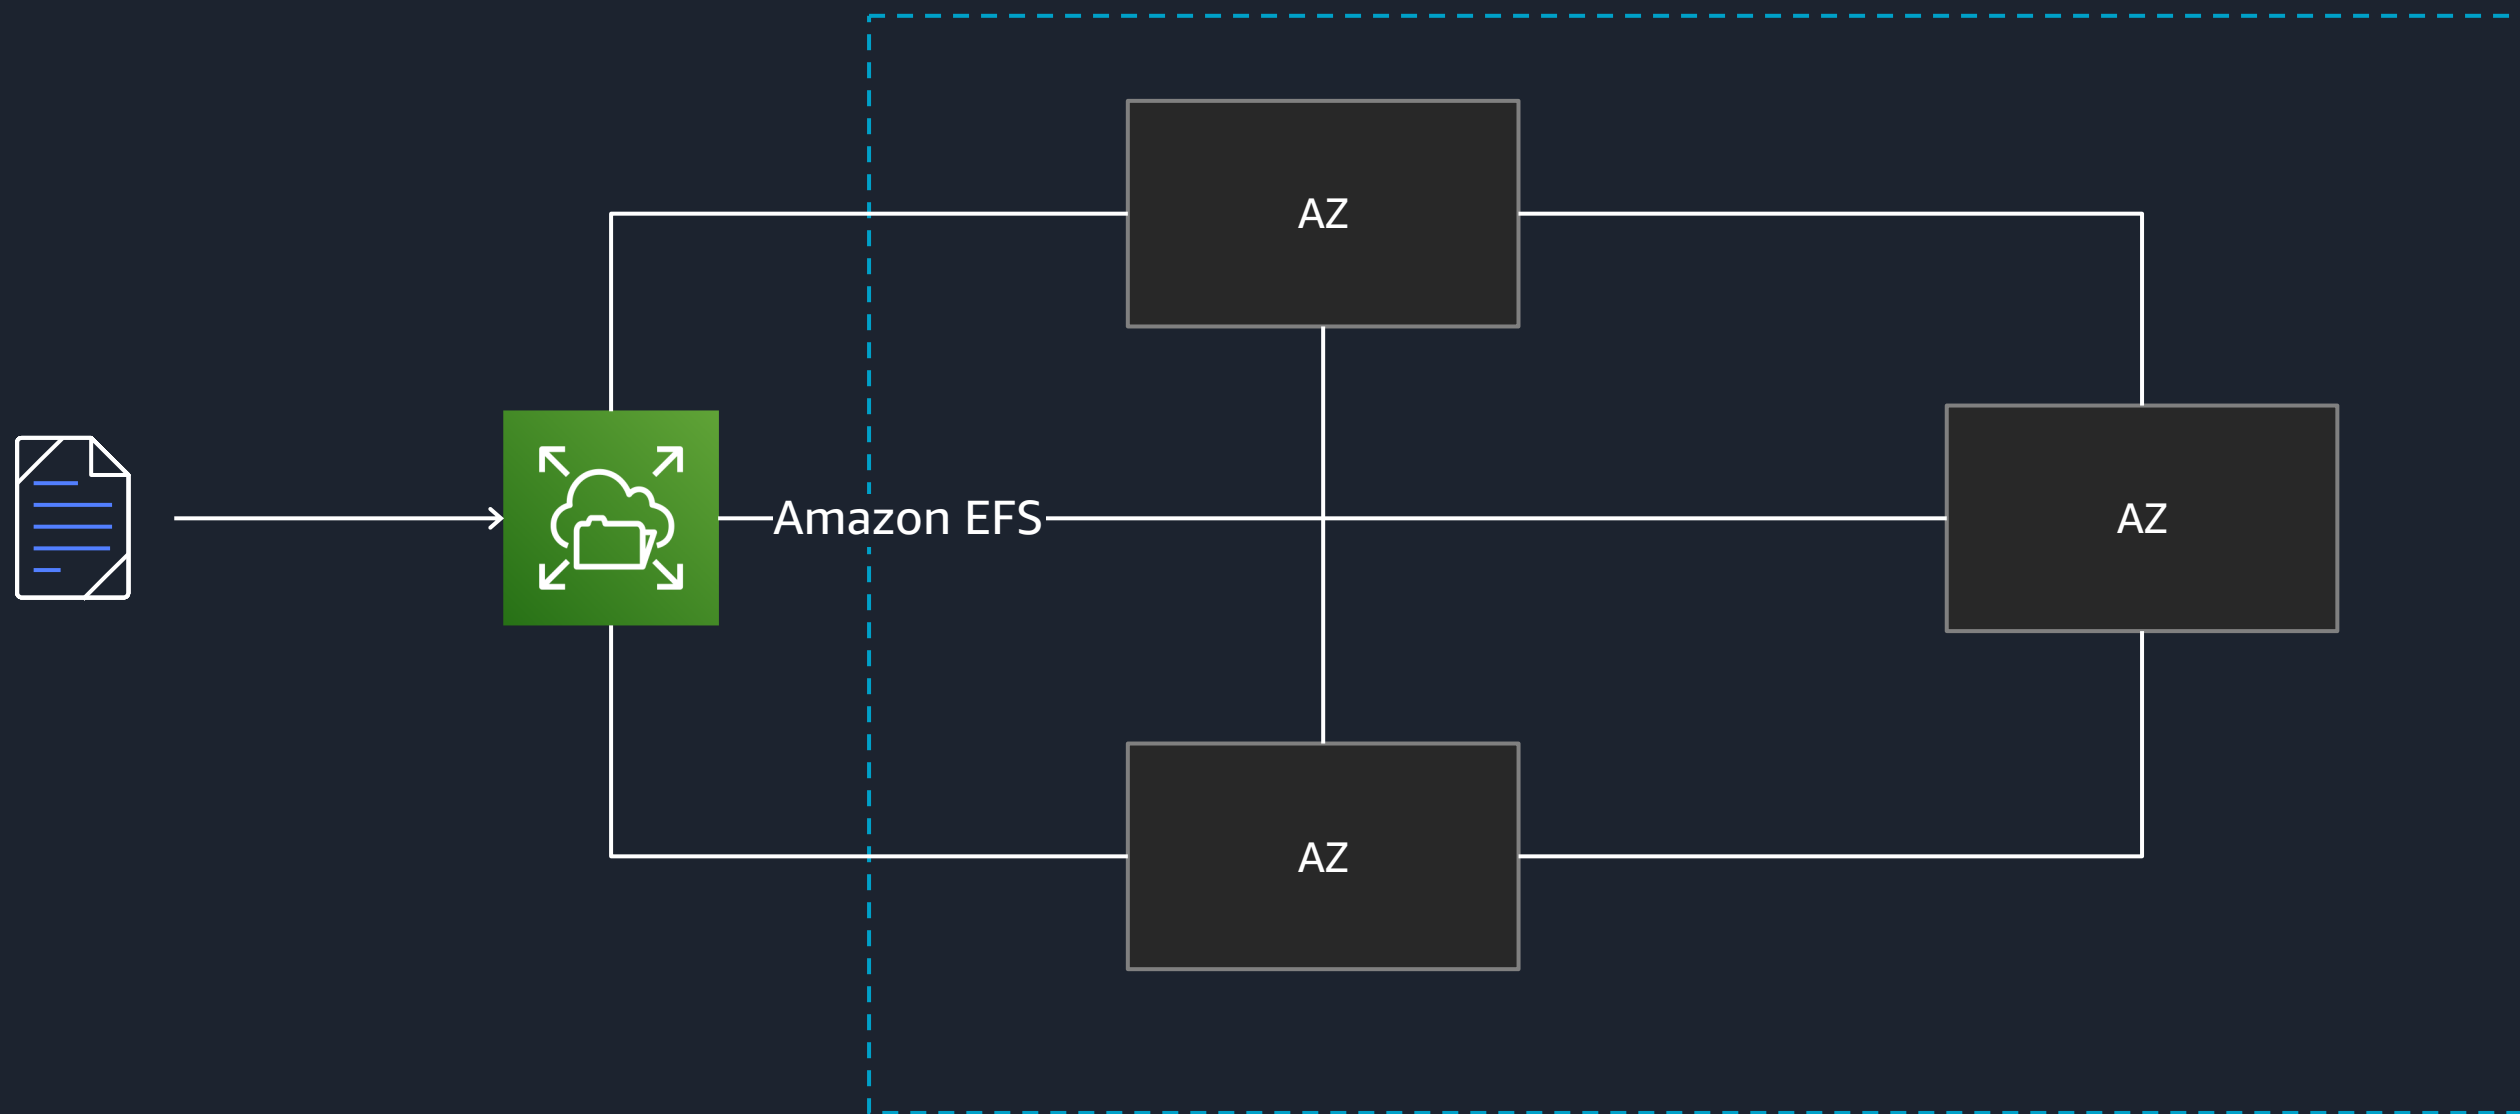
Highly integrated, shared access





Highly reliable

Benefits of 3 independent AZs in each AWS Region



99.99% Availability
SLA



Designed for
99.999999999% durability **aws** storage

Cost optimized | Amazon EFS



No minimum commitments or upfront fees




No need to provision storage



Use with Spot Instances



Automatic lifecycle management to lower cost storage



Amazon EFS

Pay-as-you-go, effective storage cost of
\$0.08/GB-month^{*,+}



Standard storage class

General-purpose file storage
\$0.30/GB-month*



Infrequent access storage class

Cost-optimized for files not accessed every day
\$0.025/GB-month* for storage
\$0.01/GB* for access

Amazon EFS Infrequent Access


Amazon EFS IA storage class for infrequently accessed files for \$0.025/GB per month*



No changes to existing
applications using
Amazon EFS



Cost
savings up
to 92%

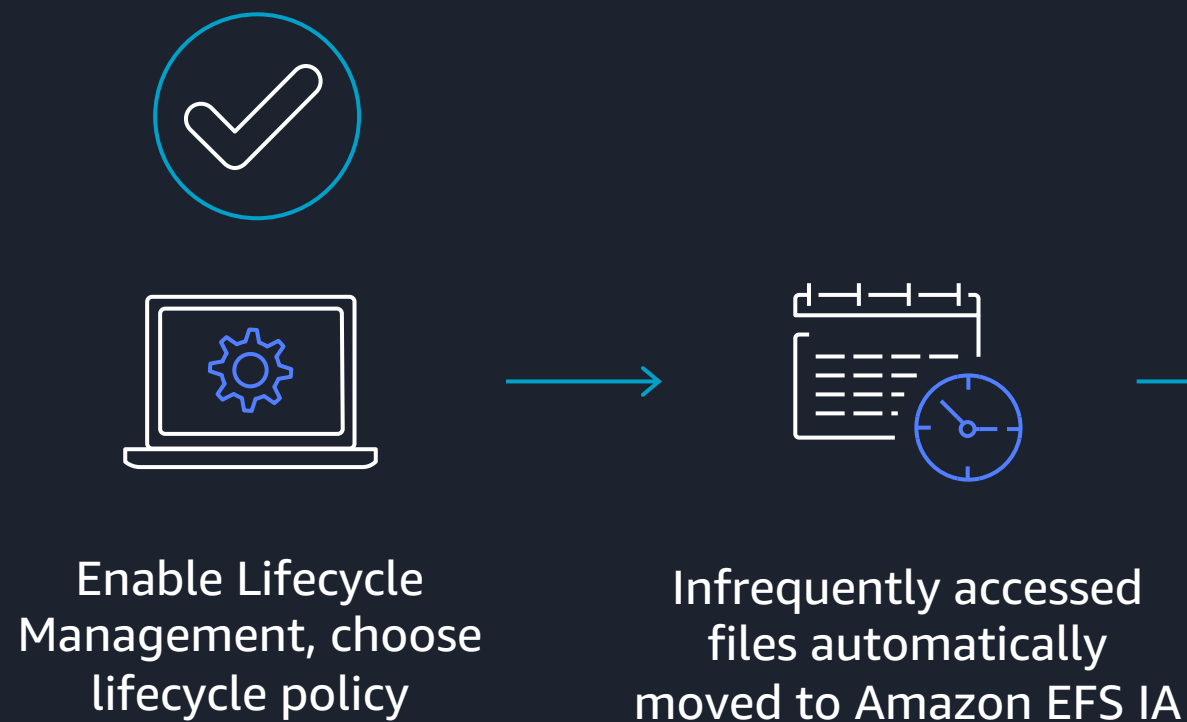


Automated
lifecycle
management

* Pricing in the US East (N. Virginia) region

Enabling EFS Lifecycle Management

Lifecycle policies can be configured to 7, 14, 30, 60, or 90 days since last access



All EFS features are supported with Amazon EFS IA

Amazon EFS performance modes



General Purpose (default)

Recommended for the
majority of workloads

EFS performance modes



General Purpose (default)

Recommended for the
majority of workloads



Max I/O

Recommended for
scale-out workloads

EFS performance modes: General purpose

What it's for

Latency-sensitive applications and general-purpose workloads

Large-scale and data-heavy applications



Advantages

Lowest latencies for file metadata operations

Virtually unlimited ability to scale out IOPS



General Purpose (default)

Recommended for the majority of workloads

Trade-offs

35k read ops/sec

Slightly higher metadata latencies

Max I/O

Recommended for scale-out workloads

When to use

Best choice for most workloads

Consider for large scale-out workloads


EFS throughput modes



Bursting Throughput (default)


Recommended for the
majority of workloads

EFS throughput modes



Bursting Throughput (default)


Recommended for the
majority of workloads



Provisioned Throughput

Recommended for
higher throughput-to-
storage ratio workloads

EFS throughput modes



Bursting Throughput (default)

Recommended for the majority of workloads

What it's for

Varying throughput workloads	Increased, more consistent throughput workloads
------------------------------	---

Advantages

Auto-scaling throughput	User-defined throughput
-------------------------	-------------------------

Trade-offs

Fixed throughput-to-storage ratio	Separate throughput charge
-----------------------------------	----------------------------

When to use


Best choice for most workloads	Higher throughput-to-storage ratio
	Loading more than 2.1 TB



Provisioned Throughput

Recommended for higher throughput-to-storage ratio workloads

Provisioned Throughput




Independent throughput

Provision throughput independently of data stored



Increase

As often as you need



Switch or decrease

Once every 24+ hours

Bursting Throughput pricing

Single pricing dimension



Storage price

\$0.30/GB-month*

Includes 50 KiB/s throughput per GiB of storage

* Pricing in the US East (N. Virginia) region



Provisioned Throughput pricing

Two pricing dimensions



Storage price

\$0.30/GB-month*

Includes 50 KiB/s throughput per GiB of storage

* Pricing in the US East (N. Virginia) region



Provisioned Throughput pricing

Two pricing dimensions



Storage price

\$0.30/GB-month*

Includes 50 KiB/s throughput per GiB of storage




Throughput price

\$6.00/MiBps-month*

Above 50 KiB/s throughput per GiB of storage

* Pricing in the US East (N. Virginia) region

Security and compliance



Control network traffic

Using Amazon VPC security groups and network ACLs



Control file and directory access

Using POSIX permissions

control client identity/mount permissions with access points



Control administrative & application access

Using AWS IAM, action-level and resource-level permissions, identity-based policies

Manage application access with EFS Access Points



Encrypt data

At rest and in transit



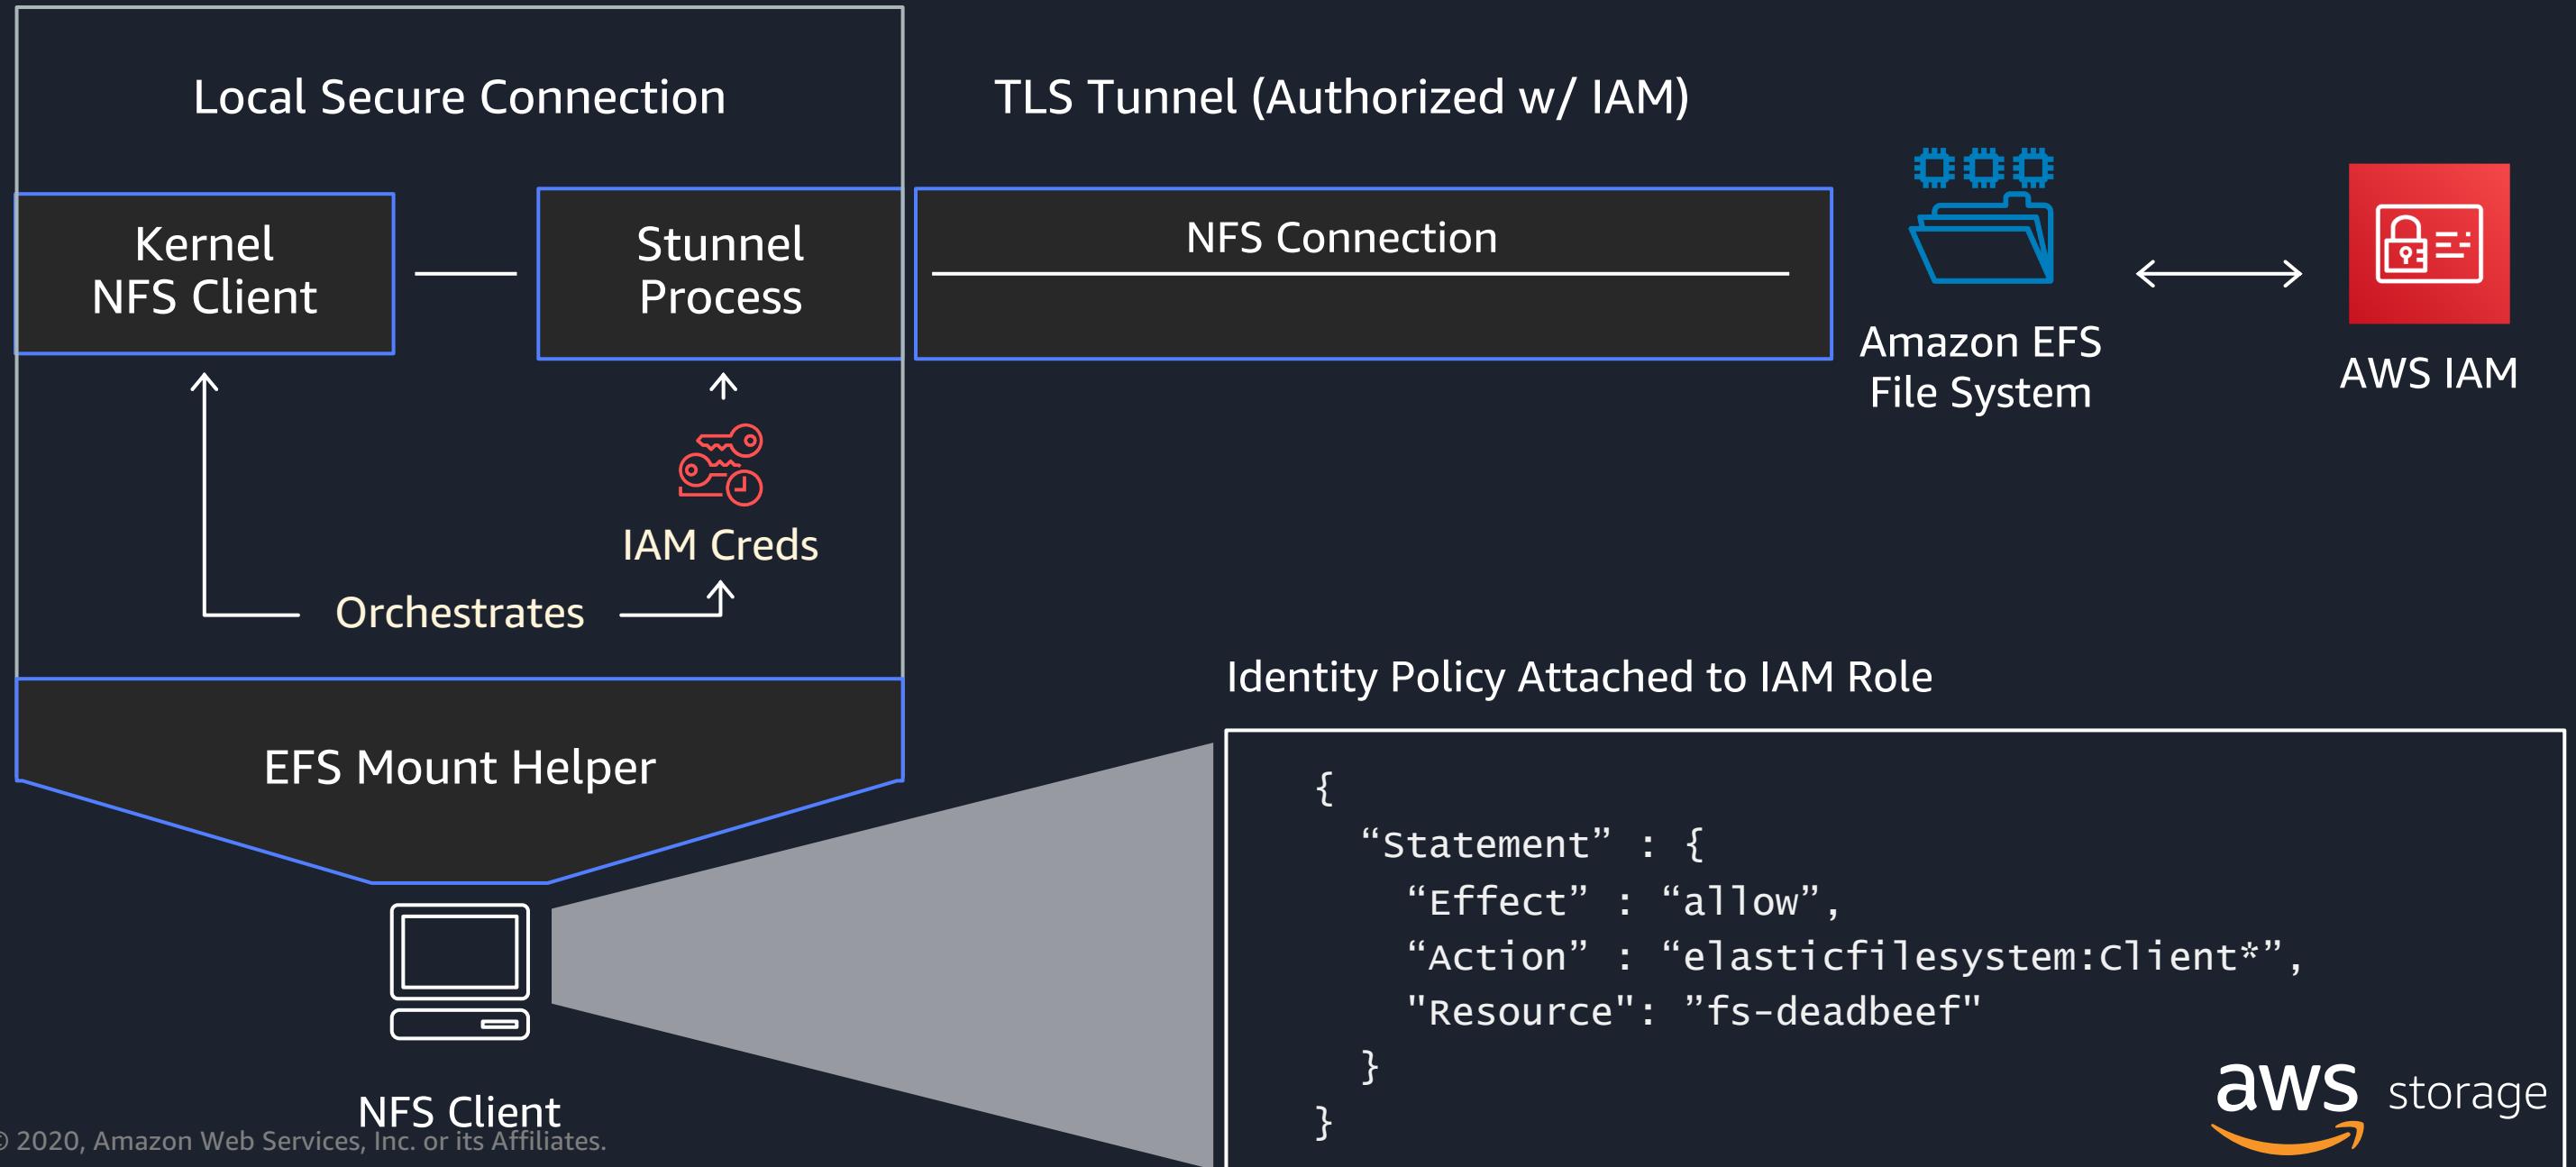
Achieve compliance

HIPAA
GDPR
PCI-DSS
SOC
ISO

FedRAMP
aws storage

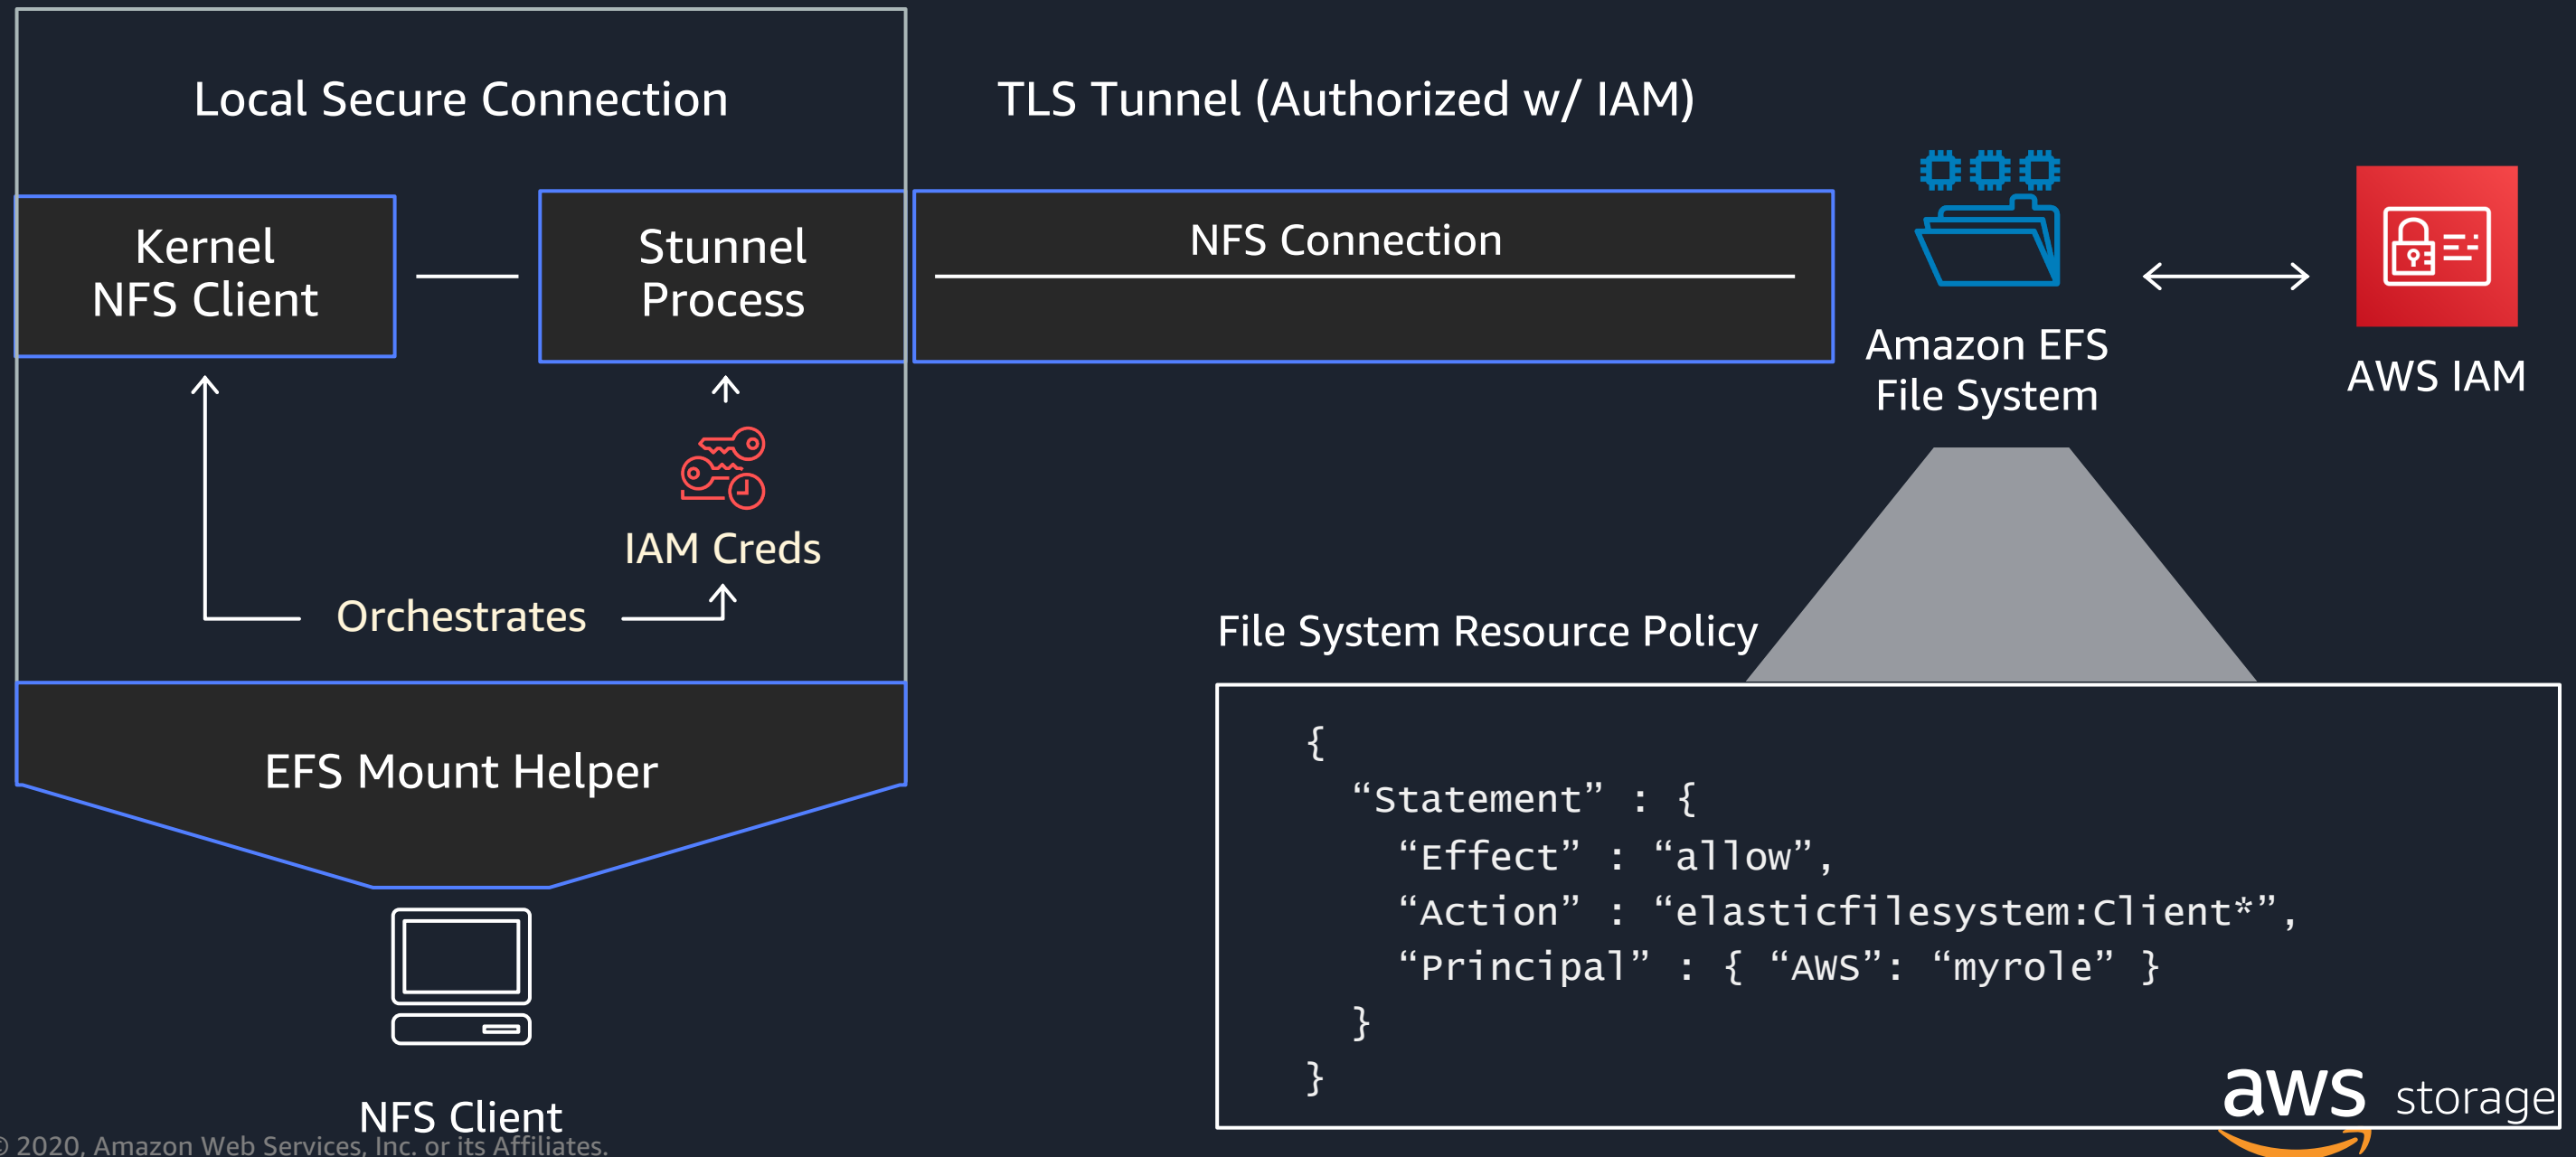
Amazon EFS | IAM authorization for NFS clients

Restricting EFS access using IAM Identity Policy



Amazon EFS | IAM authorization for NFS clients

Restricting EFS access using IAM Resource Policy



Application-specific Access with EFS Access Points

Creates App-specific Directory & Permissions

No EC2 instance required!

Apps only see data they need

Enforces File System Identity

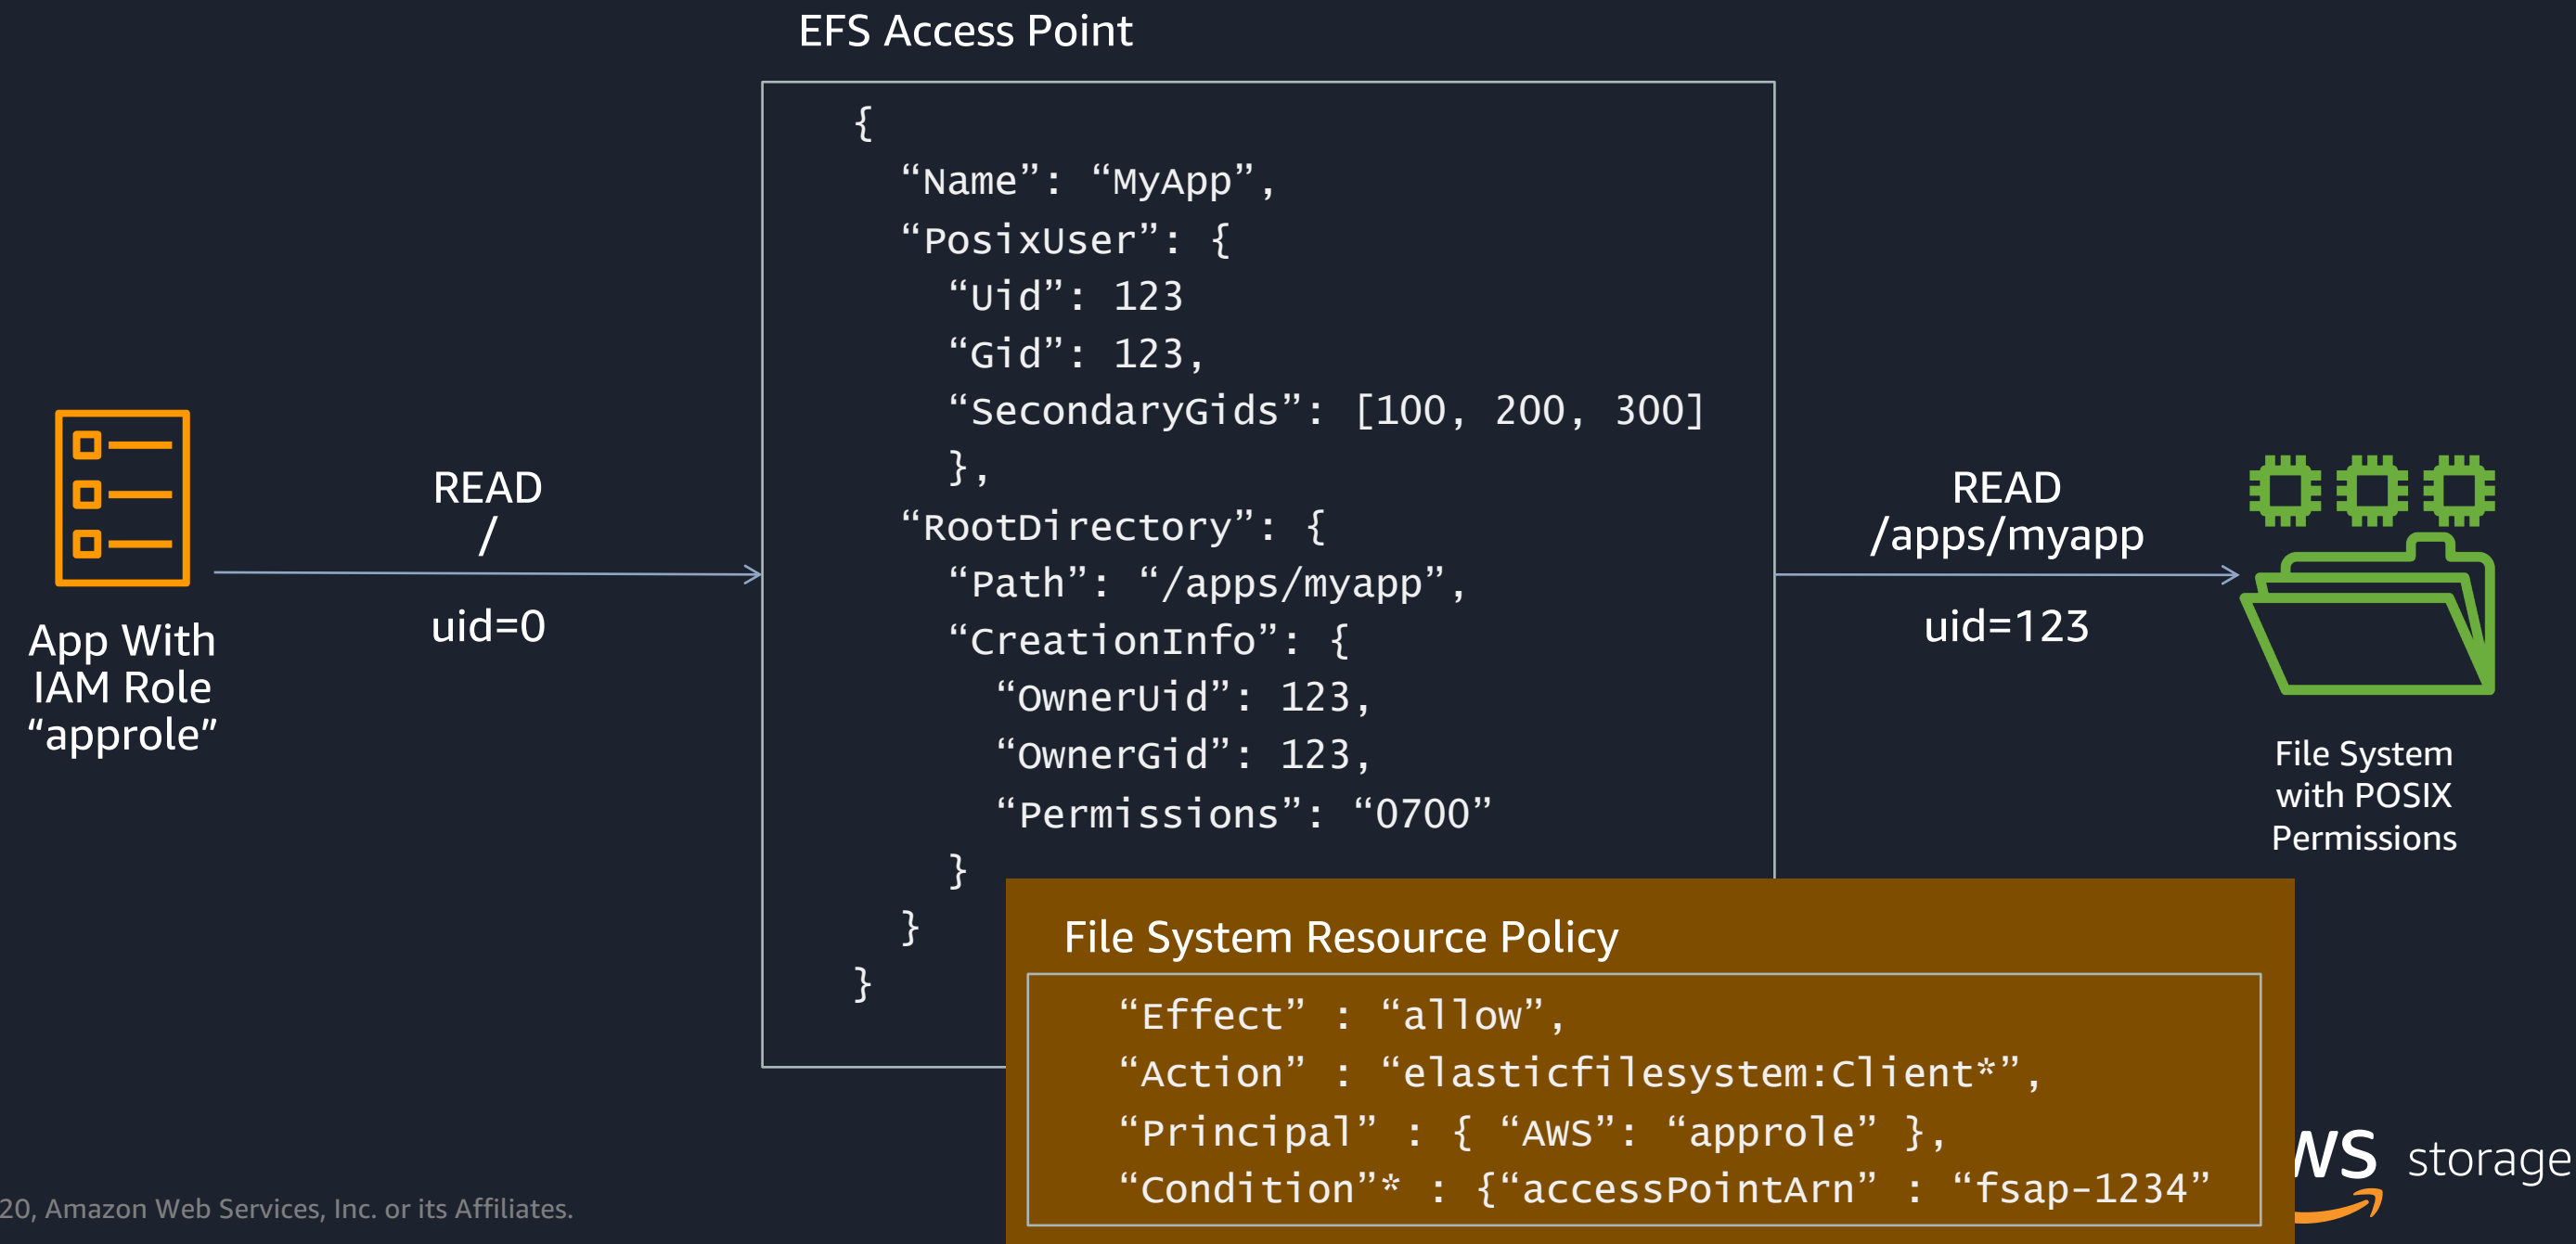
Root containers can't escalate access

Arbitrary users aren't locked out

EFS Access Point

```
{
  "Name": "MyApp",
  "FileSystemId": "string",
  "PosixUser": {
    "Uid": 123
    "Gid": 123,
    "SecondaryGids": [100, 200, 300]
  },
  "RootDirectory": {
    "Path": "/apps/myapp",
    "CreationInfo": {
      "OwnerUid": 123,
      "OwnerGid": 123,
      "Permissions": "0700"
    }
  }
}
```

How EFS Access Points Work



What's new?

Since last summit, October 2019

What's new | Since Oct' 2019



Amazon EFS

October 2019

- EC2 Launch Instance Wizard (LIW)
- AWS Private Link

December 2019

- Service-Linked Roles

February 2020

- GovCloud (US-East)
- EC2 LIW UX

April 2020

- 400% Read IOPS increase (up to 35K)
- Integration with ECS and Fargate
- 99.99% SLA

Today

November 2019

- EFS IA 7-day LM policy
- Middle East (Bahrain)
- EU (Stockholm)
- South America (Sao Paulo)
- Asia Pacific (Hong Kong)

January 2020

- AWS IAM Authorization for NFS Clients
- EFS Access Points
- Item-Level Recovery (using AWS Backup)
- China (Ningxia, Beijing)

March 2020

- AWS Budgets*

May 2020

- CPT, MXP regions

AWS integrations | Amazon EFS



Amazon EFS




Amazon
VPC




AWS
IAM




AWS
KMS




Amazon
CloudWatch




AWS
CloudTrail




AWS
CloudFormation




AWS Direct
Connect




AWS VPN



AWS
DataSync



AWS
Backup



Amazon
EC2



Amazon
SageMaker



Amazon
EKS



Amazon
ECS



AWS
Fargate

Amazon EC2 launch wizard | Amazon EFS



- Mount Amazon EFS file systems directly from the Amazon EC2 launch wizard
- Select an existing file system to mount or create a new file system

aws Services Resource Groups

1. Choose AMI 2. Choose Instance Type 3. Configure Instance 4. Add Storage 5. Add Tags 6. Configure Security Group 7. Review

Step 3: Configure Instance Details

Shutdown behavior ⓘ Stop

Enable termination protection ⓘ ☐ Protect against accidental termination

Monitoring ⓘ ☐ Enable CloudWatch detailed monitoring
Additional charges apply.

Tenancy ⓘ Shared - Run a shared hardware instance
Additional charges will apply for dedicated tenancy.

Elastic Inference ⓘ ☐ Add an Elastic Inference accelerator
Additional charges apply.

T2/T3 Unlimited ⓘ ☐ Enable
Additional charges may apply

File systems ⓘ

fs-9313e712	/mnt/efs/fs1	⊗
fs-7ec121ff	/mnt/efs/fs2	⊗
fs-6712e6e6	/mnt/efs/fs3	⊗

[Create new file system](#)

Add file system Add to user data

New:

Amazon ECS and AWS Fargate Support for Amazon EFS

Simple: All EFS configuration is inside the ECS task definition, and connectivity is handled behind the scenes.

Serverless: AWS Fargate tasks can now leverage shared persistent storage.

Secure: Access to file systems can be authorized by IAM, and access to data controlled by EFS Access Points.

Add volume

Name

NotebookServer

i

Volume type

EFS

i

File system ID

sourcecode | fs-8...

i

Access point ID

None

i

Root directory

/

i

Encryption in transit

☒ Enable Transit Encryption

i

EFS IAM authorization

☐ Enable IAM Authorization

i

*Required

Cancel

Add

sourcecode | fs-8239c3fa

trainingdata | fs-dc0d2aa5

System in

Create an access point for your file system in the [Amazon EFS Console](#)

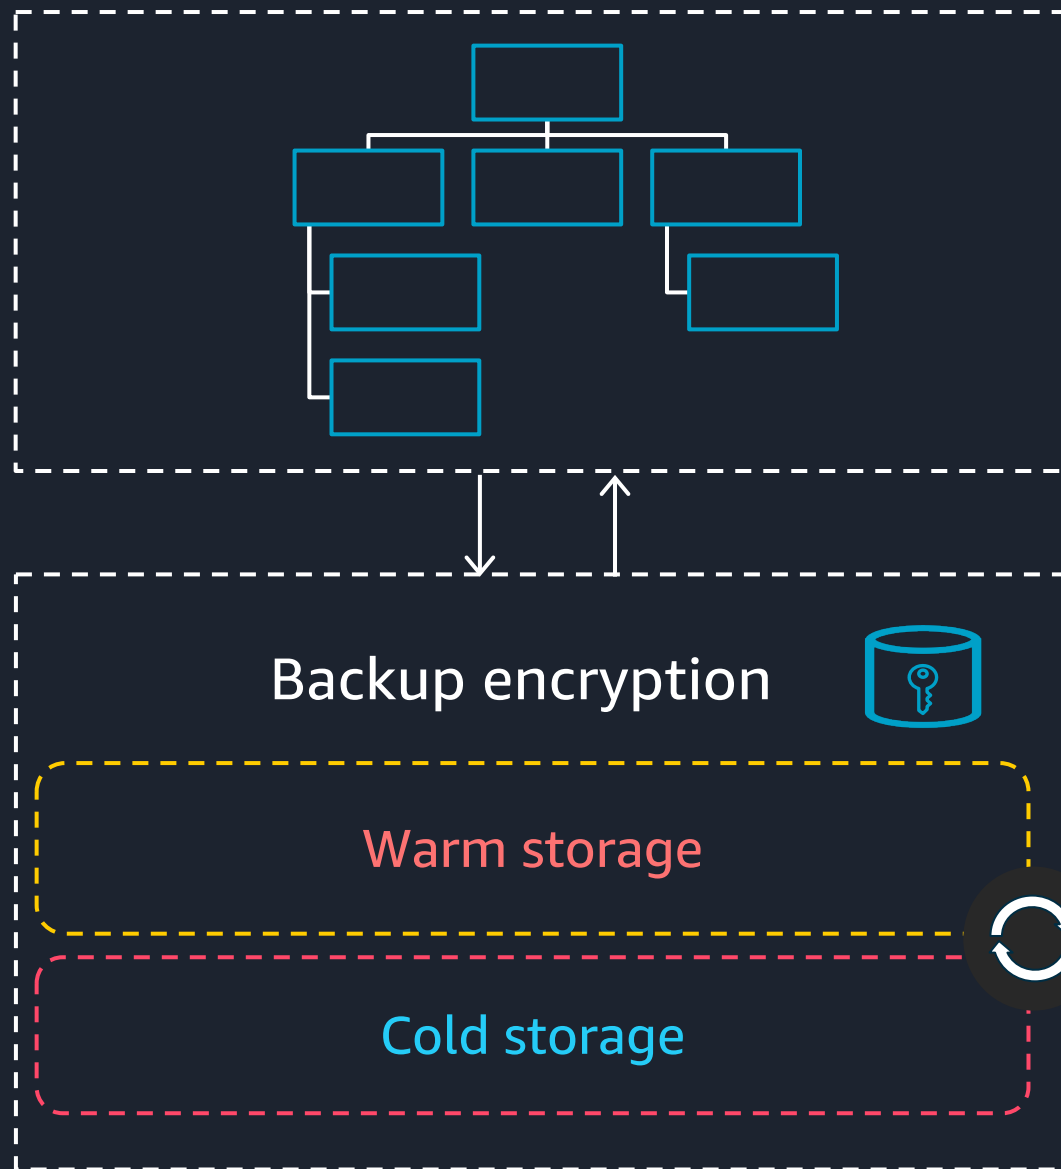
Backup for Amazon EFS



Amazon EFS

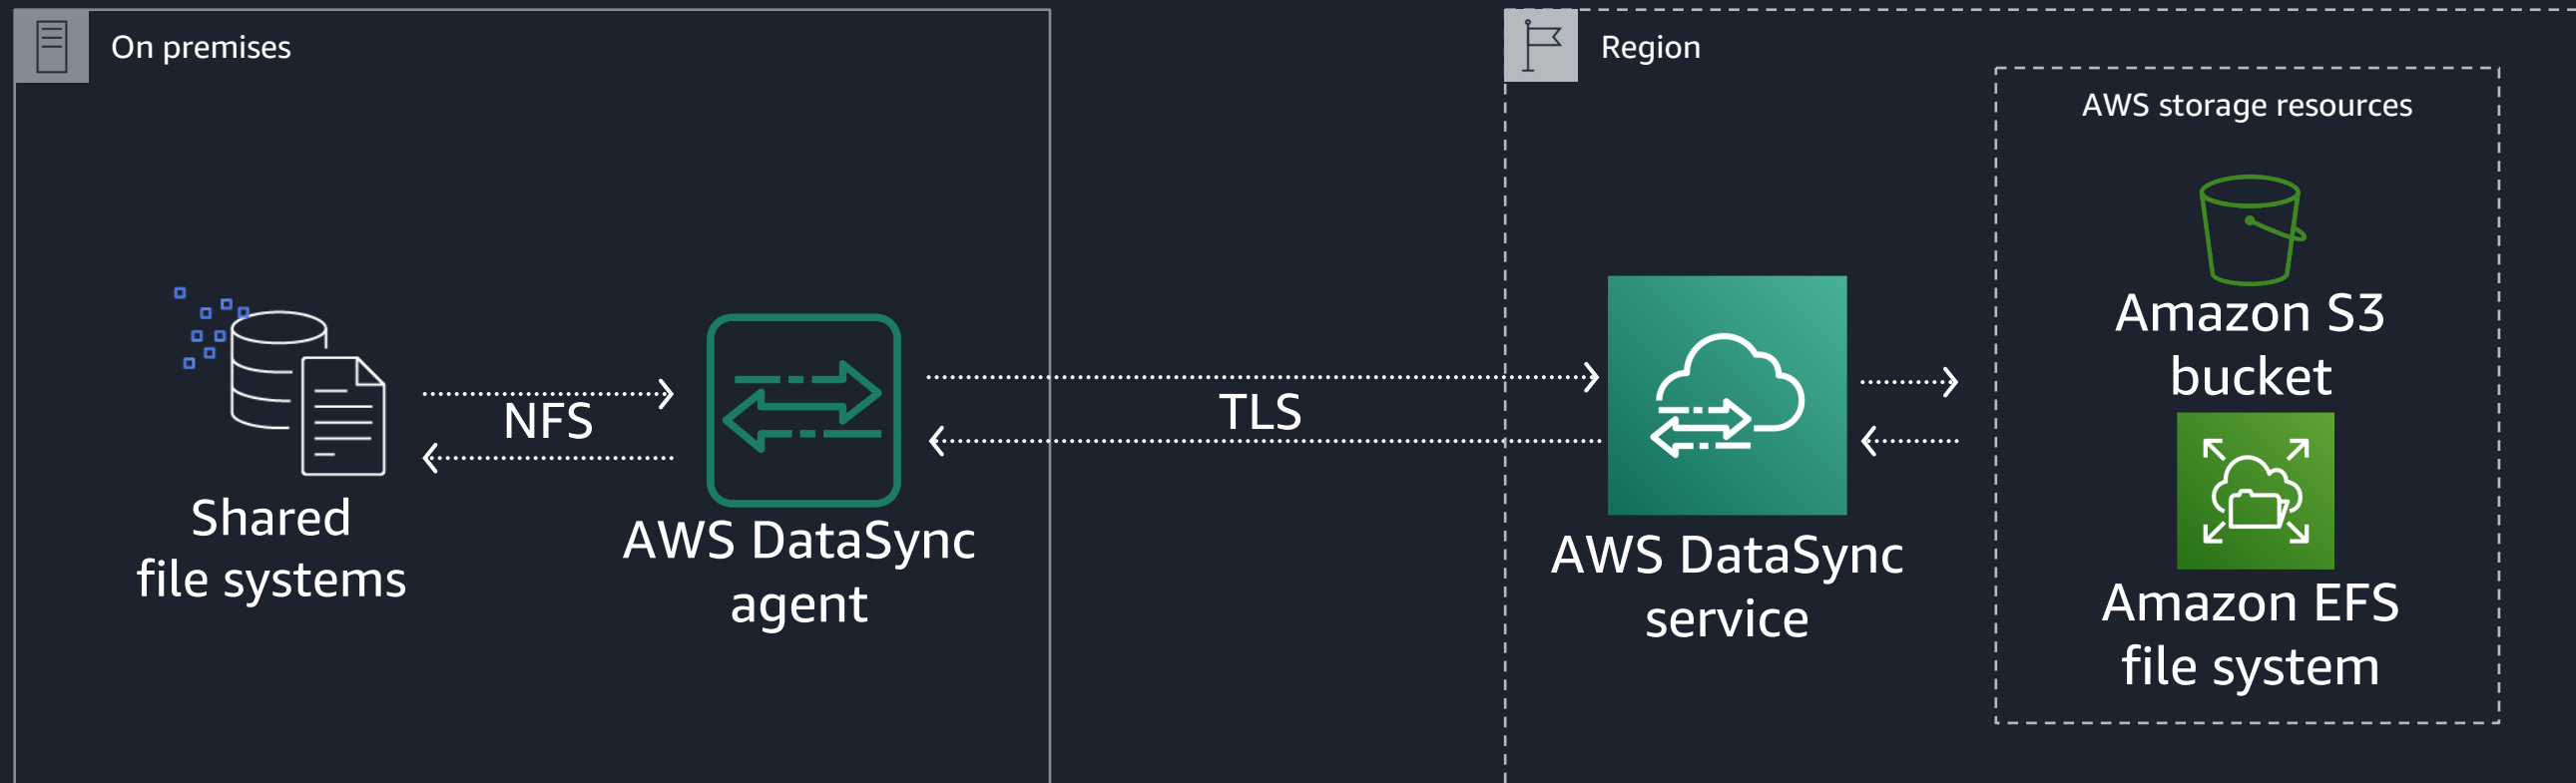


AWS Backup



- EFS file systems can be backed up and restored using AWS Backup
- AWS Backup provides automated backup scheduling and retention per user defined policy
- AWS Backup offers two classes of service backup storage with the ability to lifecycle to cold storage
- Restore individual files and directories

Migration to Amazon EFS



AWS DataSync: Online transfer service that simplifies, automates, and accelerates moving data between on-premises storage and AWS

Customer Use Cases


Use cases for Amazon EFS



Home directories
Container storage
Application test/dev



Lift and shift enterprise apps
Web serving
Content management
Database backups



Analytics
Media workflows

Metadata-intensive jobs

Scale-out jobs


Low latency and serial I/O

High throughput and parallel I/O

Demo

See the service in action!

Amazon EFS | Best practices



Test in
General Purpose
performance mode



Start with
Bursting
Throughput mode



Consider Provisioned
Throughput mode for
loading >2.1 TB




Linux kernel 4.3+



EFS mount helper
(NFSv4.1)



Large IO size
(aggregate IO)



Multiple
threads



Multiple
instances



Multiple
directories

Amazon EFS | Best practices



Enable Lifecycle Management
to automatically
save up to 92%



Enable encryption
at rest and in transit
for sensitive workloads



Create a backup
plan to further
protect your data



Monitor throughput
utilization, burst credits,
and PercentIOLimit

Getting started with Amazon EFS



Sign up for an
AWS account

Instantly get access
to the AWS Free Tier



10-minute tutorial

Learn how to use
Amazon EFS



Start building
your cloud-native
file system

Step-by-step
walkthroughs

Get started: <https://aws.amazon.com/efs/getting-started/>

Tutorials: <https://github.com/aws-samples/amazon-efs-tutorial>

Migration: <https://d1.awsstatic.com/whitepapers/aws-load-store-protect-linux.pdf>